

REGULATORY GUIDE

The Global AI Regulatory Landscape

Who's affected, what changed, and how to prepare

A year ago, AI regulation was mostly theoretical. Colorado's AI law hadn't taken effect, California's SB 1047 had just been vetoed, and the EU AI Act existed largely on paper. Over the past twelve months, that changed. Three US states built the first real accountability framework for frontier AI developers. A fourth tore up its own law and started over. The EU renegotiated its flagship regulation under industry pressure. The federal government opened a direct legal fight with the states over who gets to set the rules. None of it has settled into something stable, and most of it already applies to somebody. This guide breaks down what changed, who it reaches, and what to do about it depending on where your company sits.

QUICK REFERENCE: THE MAJOR LAWS AT A GLANCE

Law	Jurisdiction	Who's covered	Effective	Key requirement
SB 53 (Transparency in Frontier AI Act)	California	Frontier developers, \$500M+ revenue	Jan 1, 2026	Safety framework, incident reporting, whistleblower protections
RAISE Act	New York	Frontier developers, \$500M+ revenue	Dec 2025, amended Mar 2026	Same model as SB 53, \$1M/\$3M penalties
AI Safety Measures Act (AISMA)	Illinois	Large frontier developers, \$500M+ revenue	Jan 1, 2027	All of the above, plus mandatory annual third-party audits
SB 26-189 (ADMT)	Colorado	Developers and deployers of automated decision tech	Jan 1, 2027	Risk management, impact assessments, disclosures
TRAIGA (HB 149)	Texas	Developers and deployers doing business in Texas	Jan 1, 2026	Prohibits specific restricted uses (CSAM, deepfakes, self-harm)

Law	Jurisdiction	Who's covered	Effective	Key requirement
EU AI Act, Annex III (standalone high-risk)	European Union	Providers/deployers in hiring, credit, education, law enforcement	Dec 2, 2027	Risk management, documentation, conformity assessment
EU AI Act, Annex I (embedded high-risk)	European Union	AI embedded in regulated products (medical, machinery, vehicles)	Aug 2, 2028	Aligned with existing product safety regimes
EU AI Act, Article 50	European Union	Any provider of generative AI reaching EU users	Aug 2, 2026	AI-interaction disclosure, watermarking, provenance
Executive Order 14365	US federal	All US states with AI laws	Dec 11, 2025	DOJ litigation task force aimed at preempting state AI laws

WHO'S ACTUALLY AFFECTED

AI regulation splits into obligations that fall on a handful of companies and obligations that fall on nearly everyone who touches AI in a regulated way.

- **A small number of frontier developers** carry the heaviest new load. California, New York, and Illinois all set a \$500M revenue threshold, which in practice means five to ten companies: OpenAI, Anthropic, Google DeepMind, Meta, Microsoft, and a few others.
- **Employers and anyone using AI in hiring, credit, or other consequential decisions** fall under Colorado's ADMT law, NYC Local Law 144, and the EU AI Act's Annex III category. This is the category most companies actually fall into, even if they've never thought of themselves as an "AI company."
- **Anyone shipping generative AI features into the EU** has a watermarking and content-provenance obligation coming under Article 50, regardless of company size.
- **Companies building or deploying companion chatbots reachable by minors** face fast-moving state rules on crisis-response protocols and disclosure.
- **Everyone else** is affected indirectly through vendor risk, data provenance obligations, and a general legal expectation that a company can explain what its AI systems do with data.

THE US FRONTIER THREE: CALIFORNIA, NEW YORK, ILLINOIS

California SB 53 came first, signed September 2025 and in force since January 2026. It requires large frontier developers to publish safety frameworks, report critical incidents, and protect whistleblowers, and lets a company satisfy the law by meeting a comparable federal or foreign standard instead of filing twice.

New York's RAISE Act, signed December 2025 and amended in March 2026 to align with California's approach, now carries penalties of up to **\$1M** for a first violation and **\$3M** for repeat offenses, down from the \$10M/\$30M the legislature originally passed.

Illinois's AI Safety Measures Act, signed July 6, 2026 and effective January 1, 2027, uses the same \$500M threshold but adds a first-in-the-nation requirement: mandatory annual independent third-party audits, every year, for as long as a company stays covered.

Together, these three states cover roughly **40 percent** of the US AI market, and lawmakers in all three have said plainly this is a deliberate attempt to build a de facto national standard while Congress stays on the sidelines.

THE COMPREHENSIVE LAWS THAT COULDN'T HOLD THEIR SHAPE

Colorado got there first with SB 205 in 2024, but it never took effect. Delayed twice under industry pressure, it was repealed and replaced by **SB 26-189**, which narrows the focus from broad "high-risk AI" to a specific automated decision-making technology regime, effective January 2027.

Texas TRAIGA, in force since January 2026, took the narrower route from the start: rather than broad risk assessments, it prohibits specific uses, including AI-generated CSAM, unlawful deepfakes, and systems designed to encourage self-harm.

Comprehensive, high-risk AI acts modeled on Colorado's original approach have mostly lost ground to narrower, harm-specific rules and to the frontier-developer model above.

THE EU AI ACT AND THE DIGITAL OMNIBUS: NEW TIMELINE, SAME RULES

The EU AI Act is still the broadest AI framework in the world, and it's still in force. What changed this year is the timeline, and in a few places, the substance. The **European Commission** published the **Digital Omnibus on AI** on November 19, 2025, proposing to defer the most demanding parts of the Act. Parliament and Council reached political agreement on May 7, 2026, formally adopted through June.

Where the timeline landed

- High-risk obligations, standalone systems (Annex III): moved to **December 2, 2027**
- High-risk obligations, embedded systems (Annex I): moved to **August 2, 2028**
- General transparency (Article 50): unaffected, still **August 2, 2026**
- Watermarking (Article 50(2)): pushed to **December 2, 2026** for systems already on the market
- New prohibition on non-consensual intimate imagery and CSAM generation: **December 2, 2026**
- Foundation model obligations (Articles 51-55): in force since August 2025, untouched

The underlying obligations survived intact; only the calendar moved. Companies treating the extra sixteen months as a pause are setting up for a compressed scramble in late 2027. The companies in the strongest position are using this window to inventory every AI system in use, classify it against Annex III and Annex I, and map what data each one touches.

THE NARROWER BILLS THAT ADDED UP FAST

Away from the headline frameworks, a much larger volume of narrower legislation reshaped specific corners of the landscape. Companion chatbots and mental health became their own category, largely in response to reported harms to minors: **Illinois** barred unlicensed AI from providing psychotherapy, **California's Companion Chatbots Act** now requires crisis-response protocols, and **New York** has similar legislation in progress. Content provenance rules are catching up with generative media on both sides of the Atlantic, and employment AI oversight kept expanding around **NYC's Local Law 144** and its bias-audit requirements for hiring tools.

THE US FEDERAL COLLISION COURSE

While states legislated, the federal government built toward a fight over whether any of it can stand. **Executive Order 14365**, signed December 11, 2025, directs the DOJ to stand up an AI Litigation Task Force to sue states over "onerous" AI laws, ties broadband funding to state compliance, and names Colorado's law directly as a target. It carves out child safety, AI infrastructure, and government procurement from any preemption push. An executive order cannot preempt state law on its own: only Congress or the courts can, and neither has yet. The advice from every firm tracking this is the same: keep complying with the state laws that apply to you now.

HOW TO PREPARE, DEPENDING ON WHERE YOU SIT

If you're a large frontier developer (CA, NY, or IL thresholds)

A published, annually updated safety framework, incident reporting infrastructure with tight timelines (72 hours in most cases), whistleblower protection processes, and, starting with Illinois, standing readiness for an independent third-party audit every year.

If you're deploying AI in hiring, lending, or other consequential decisions

Colorado's ADMT law, NYC Local Law 144, and the EU AI Act's Annex III category all converge on the same requirement: know what data feeds the decision, document it, and be able to produce a bias audit or impact assessment on demand.

If you ship generative AI features into the EU

Watermarking and disclosure infrastructure needs to be operational by December 2, 2026 for systems already on the market. Seven months is a tight window for UI labeling, metadata embedding, and detection capability.

If none of the above apply directly

Inventory what AI systems you use, built or bought, and know what data each one touches. The only durable strategy is data-level readiness that doesn't depend on which jurisdiction's paperwork wins out.

4 STEPS TO GET AHEAD OF THE AI REGULATORY WAVE

Every obligation above, regardless of jurisdiction, comes back to the same operational question: what data does this AI system touch, and can you prove it. That question increasingly extends past models to the agents built on top of them. This is where a data control plane matters: a single layer that governs data itself, together with every model and agent that touches it, rather than bolting governance onto each tool after the fact. BigID operates as that control plane, built around contextual risk, which means seeing the full picture across data and agents and then acting on it.

- 1 Find every AI system and agent in use, and the data each one can reach.** Compliance teams can't classify or document a system they don't know is running. BigID discovers AI tools, models, and agents, and automatically understands the data they touch.
- 2 Know what data trained or feeds each model.** Auditors, notified bodies, and internal risk committees all ask where a model's data came from and what it can reach now. BigID maps data lineage and model-to-data relationships automatically.
- 3 Classify sensitive data before an AI system or agent touches it.** Bias assessments and risk reviews depend on knowing which data is sensitive in the first place. BigID classifies data at the source across structured and unstructured systems.

- 4 Govern and monitor access, then keep the evidence.** Whistleblower protections, reporting windows, and annual audits all assume you can show who and what could reach a model and its data at any given time. BigID monitors access continuously and retains that record.

The regulatory theories differ by state and by continent. The operational answer, in nearly every case, starts in the same place: knowing your data well enough to show your work, whenever a regulator asks.

See it in action

The fastest way to understand how this applies to your own environment is to see it against your own data. Take a self-guided tour or book a live demo to walk through how discovery, lineage, classification, and access governance map to the specific laws your company falls under.

CUSTOMER VALIDATION

Proof in Action

"BigID gives me better visibility into sensitive data, helps prioritize security protections, reduces attack surfaces, strengthens compliance, and increases operational efficiency overall — a strategic pillar of our AI-First Cybersecurity Transformation"

- CISO, Global Healthcare Company

"We needed to automate processes and data management across systems for the data of a few million customers across a lot of systems and data. BigID was the one solution that did this in the most efficient and sophisticated way - and had more use cases we could add on moving forward."

- Chief Privacy Officer, Global Telecoms Company

ANALYST RECOGNITION

Validated by Industry Leaders

- 2025 Company of the Year for AI Governance - Frost & Sullivan
- Named a Leader across all DSPM research including Omdia DSPM Universe, CB Insights for DSPM, TAG, GigaOm, Frost & Sullivan
- Named a Leader in all Privacy Management evals including Forrester Privacy Management Wave & IDC Privacy Compliance MarketScape
- BigID named a Leader across 4 GigaOm evals: DSPM, Data Security Platforms, Unstructured Data Management, Data Access Governance
- Named in Gartner's 2025 DLP Market Guide
- Named in Gartner's 2025 AI TRiSM Market Guide
- Represented 30+ times in Gartner's 2025 Hype Cycles for Security, Privacy, AISPM

FREE RISK ASSESSMENT

Discover Your Data & AI Risk

An agentless, cloud-native assessment to enable security, privacy, and compliance: pinpoint where sensitive data and AI risks live - and how to fix them.

What You Get:

- Complementary 2 week engagement on real data to gain visibility on risk across data and AI
- Executive-ready report identifying high-risk data and AI across cloud environments with next-step recommendations

NEXT STEPS

Ready to Connect the Dots in Data & AI?

Schedule your free risk assessment or ask for a deeper demo across DSPM, DLP, AI SPM, Privacy Automation, AI Risk, and Data Lifecycle Management

► www.bigid.com/demo

► www.bigid.com/risk-assessment

► info@bigid.com



BigID Next

"The features they offer are the moat that they have built around the product. This is what makes them the market leader in Data Privacy and Protection."



IT Security & Risk Management Associate - Banking

Gartner
Peer Insights.

Gartner® and Peer Insights™ are trademarks of Gartner, Inc. and/or its affiliates. All rights reserved. Gartner Peer Insights content consists of the opinions of individual end users based on their own experiences, and should not be construed as statements of fact, nor do they represent the views of Gartner or its affiliates. Gartner does not endorse any vendor, product or service depicted in this content nor makes any warranties, expressed or implied, with respect to this content, about its accuracy or completeness, including any warranties of merchantability or fitness for a particular purpose.



BigID Next

"Exceptionally powerful and easy to use tool that integrated smoothly into our environment and delivered value from day one. The BigID team is fabulous to work with."



Senior Vice President & Chief Data Officer - Banking

Gartner
Peer Insights.

Gartner® and Peer Insights™ are trademarks of Gartner, Inc. and/or its affiliates. All rights reserved. Gartner Peer Insights content consists of the opinions of individual end users based on their own experiences, and should not be construed as statements of fact, nor do they represent the views of Gartner or its affiliates. Gartner does not endorse any vendor, product or service depicted in this content nor makes any warranties, expressed or implied, with respect to this content, about its accuracy or completeness, including any warranties of merchantability or fitness for a particular purpose.